

SULAMAN KHAN

Lead Backend Engineer / Architect | TypeScript • Node.js • GraphQL • PostgreSQL

Lahore, Pakistan • Remote (UTC+5) • +92 316 7852626 • se.sulaman.khan@gmail.com
github.com/sulaman0 • linkedin.com/in/suleman-khan-eng • stackoverflow.com/users/8344330
Full overlap with Gulf hours • 3–4 hours with UK/EU • US mornings covered

PROFESSIONAL SUMMARY

Backend engineer with 8 years of production experience, operating at lead and architect level: I own systems from schema to production and make the architecture calls rather than waiting for them. I have taken a 1M-user service through a runtime migration I proposed and executed, built and owned a full product single-handedly from admin dashboard to deployed API to billing, and led engineering teams to CMMI 2.0 certification. My current work is production AI systems, so an AI-first workflow is how I already build: I use Claude, Copilot and Cursor daily and review every generated line, because the ownership does not transfer to the tool.

CORE SKILLS

Architecture: System design, service decomposition, microservices, event-driven architecture, API contract design, technical direction setting, scalability and maintainability trade-offs, migration planning

Backend: TypeScript, Node.js, Express, Go, Python, Rust, PHP/Laravel, GraphQL, REST, gRPC, Protocol Buffers, message queues, concurrency

Data: PostgreSQL, MySQL, Redis, Elasticsearch, Prisma, schema design, indexing strategy, query profiling and optimization, data modelling

Reliability & quality: Datadog monitoring and alerting, structured logging, load testing, retry with exponential backoff, graceful degradation, testing strategy, code review standards, CI/CD

Cloud & DevOps: AWS (Lambda, API Gateway), DigitalOcean, Docker, Kubernetes, Nginx, Cloudflare, GitHub Actions

AI-first delivery: Claude, GitHub Copilot, Cursor in daily workflow with full review ownership; professionally builds LLM systems — LangChain, LangGraph, RAG pipelines, agent orchestration, evaluation harnesses, structured outputs

Leadership: Led projects to CMMI 2.0 certification, code review and engineering standards, mentoring, stakeholder communication, translating ambiguous requirements into technical solutions

PROFESSIONAL EXPERIENCE

AI Engineer

Jan 2024 – Present

Vaival Technologies • Client: Vulcan Forged (US) • Remote

- **Owned the AI service architecture:** Designed and shipped LLM systems as async services with token streaming, per-request cost and latency observability, provider failover and retry with exponential backoff. Chose the abstractions, not just the implementation.
- **Made the isolation call:** Enforced multi-tenant isolation inside the retrieval query rather than post-filtering results, because post-filtering still reads other tenants' rows out of the index. Backed it with security tests seeding two tenants with deliberately similar data and asserting zero cross-tenant results.
- **Set the quality bar:** Built an evaluation harness measuring retrieval hit-rate and answer faithfulness, run on every change to chunking, embeddings or prompts, so releases were decided by numbers rather than judgement calls.
- **AI-first in practice:** Use Claude, Copilot and Cursor daily and review every generated line before it ships. Building LLM systems professionally gives me a concrete sense of where these tools are reliable and where their output has to be discarded.

Backend & Blockchain Engineer

Dec 2021 – Mar 2024

Vaival Technologies • Client: Vulcan Forged (US) • Remote

- **Proposed and executed a runtime migration:** A Node.js service serving 1M users had memory climbing under load, mitigated only by scheduled restarts. I profiled it, contained the user-facing symptom, then made the case for migrating to Go rather than patching the same class of bug repeatedly — and delivered the migration.

- **Owned services end to end:** Designed and built the block explorer from indexing blocks off the node through storage, API and frontend, plus DEX, NFT marketplace, bridge and swap modules, all serving real-time on-chain data over GraphQL across a microservices environment.
- **Database and throughput ownership:** Optimized an API to 1M+ requests through server-resource management, caching and query tuning. Designed the indexing and storage model for historical chain queries, balancing storage cost against query latency.
- **Engineering in an unforgiving environment:** Transaction signing, DeFi protocol integration and smart-contract interaction on a public network where state is permanent and a bad parameter cannot be quietly rolled back. Every change went through load testing before release.

Full Stack Engineer (Web & Mobile)

Jun 2018 – Dec 2021

B4USoft Pvt. Ltd. · Lahore, Pakistan

- **Technical leadership:** Led projects to CMMI 2.0 certification, establishing review practices and engineering standards across the team, and mentored engineers. Raised an app store rating from 4.3 to 4.9 and cut startup time to 10 seconds.
- **Fintech:** Built secure transaction, authentication, fraud-prevention and data-security features. The launched application processed \$1M+ in transactions in its first quarter.
- **Scale and cost decisions:** Optimized a CMS to 100,000+ daily active users; moved workloads to serverless on AWS Lambda with Docker deployment and Datadog monitoring, cutting operational costs 40% while improving scalability.

OWNED PRODUCTS

- **Looqer — AI Shopify App, sole owner, architecture through production:** React and Polaris admin dashboard, Node.js/Express/Prisma API, vanilla-JS widget runtime embedded in third-party storefronts, and a Shopify theme extension. Deployed on DigitalOcean with Neon PostgreSQL, Nginx, SSL and Cloudflare, with tiered subscription billing across four plans. Every architectural decision was mine.
- **ZATCA Phase 2 E-Invoicing:** Saudi tax-compliance implementation — Fatoora integration, compliant XML generation, certificate signing. Regulated engineering where the specification admits no approximation. Open-sourced; 13 stars, 7 forks, in use by other teams.
- **Shopify Apps:** Cartly, Sales Pop-up and Wheelify, used by thousands of merchants, shipped inside another platform's API limits, billing rules and review process.

SELECTED ACHIEVEMENTS

- Migrated a 1M-user Node.js service to Go — a decision I proposed, argued for and delivered.
- Optimized an API to 1M+ requests; scaled a CMS to 100,000+ daily active users.
- Fintech application processed \$1M+ in transactions within its first quarter post-launch.
- Cut operational costs 40% through serverless architecture while improving scalability.
- Led projects to CMMI 2.0 certification; app store rating 4.3 to 4.9.

EDUCATION

B.S. Information Technology

2014 – 2018

University of the Punjab, Lahore, Pakistan

- Applied Concurrency in Go — goroutines, channels, concurrent systems design (2023)